



In [1]: `!pip install -q hvplot`

In [2]: `import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import hvplot.pandas
from scipy import stats

%matplotlib inline
sns.set_style("whitegrid")
plt.style.use("fivethirtyeight")`

In [3]: `data = pd.read_csv("heart.csv")
data.head()`

Out[3]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach
0	63	1	3	145	233	1	0	150
1	37	1	2	130	250	0	1	187
2	41	0	1	130	204	0	0	172
3	56	1	1	120	236	0	1	178
4	57	0	0	120	354	0	1	163

In [4]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

In [5]: `data.shape`

Out[5]: (303, 14)

In [6]: `pd.set_option("display.float", "{:.2f}".format)
data.describe()`

Out[6]:

	age	sex	cp	trestbps	chol	fbs
count	303.00	303.00	303.00	303.00	303.00	303.00
mean	54.37	0.68	0.97	131.62	246.26	0.15
std	9.08	0.47	1.03	17.54	51.83	0.36
min	29.00	0.00	0.00	94.00	126.00	0.00
25%	47.50	0.00	0.00	120.00	211.00	0.00
50%	55.00	1.00	1.00	130.00	240.00	0.00
75%	61.00	1.00	2.00	140.00	274.50	0.00
max	77.00	1.00	3.00	200.00	564.00	1.00

In [7]: `data.target.value_counts()`

Out[7]: 1 165
0 138
Name: target, dtype: int64

In [8]: `data.target.value_counts().hvplot.bar(
 title="Heart Disease Count", xlabel='Heart Disease', ylabel='Count',
 width=500, height=350
)`

Out[8]:

In [9]: *# Checking for missing values*
`data.isna().sum()`

Out[9]: age 0
sex 0
cp 0
trestbps 0
chol 0
fbs 0
restecg 0
thalach 0
exang 0
oldpeak 0
slope 0
ca 0
thal 0
target 0
dtype: int64

In [10]: `categorical_val = []
continuous_val = []`

```
continuous_val = []  
for column in data.columns:  
    if len(data[column].unique()) <= 10:  
        categorical_val.append(column)  
    else:  
        continuous_val.append(column)
```

In [11]: categorical_val

Out[11]: ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal', 'target']

```
In [12]: have_disease = data.loc[data['target']==1, 'sex'].value_counts().hvplot.bar(alpha=0.4)  
no_disease = data.loc[data['target']==0, 'sex'].value_counts().hvplot.bar(alpha=0.4)  
  
(no_disease * have_disease).opts(  
    title="Heart Disease by Sex", xlabel='Sex', ylabel='Count',  
    width=500, height=450, legend_cols=2, legend_position='top_right'  
)
```

Out[12]:

```
In [13]: have_disease = data.loc[data['target']==1, 'cp'].value_counts().hvplot.bar(alpha=0.4)  
no_disease = data.loc[data['target']==0, 'cp'].value_counts().hvplot.bar(alpha=0.4)  
  
(no_disease * have_disease).opts(  
    title="Heart Disease by Chest Pain Type",  
    xlabel='Chest Pain Type', ylabel='Count',  
    width=500, height=450, legend_cols=2, legend_position='top_right'  
)
```

Out[13]:

```
In [14]: have_disease = data.loc[data['target']==1, 'fbs'].value_counts().hvplot.bar(alpha=0.4)  
no_disease = data.loc[data['target']==0, 'fbs'].value_counts().hvplot.bar(alpha=0.4)  
  
(no_disease * have_disease).opts(  
    title="Heart Disease by fasting blood sugar", xlabel='fasting blood sugar > 120 mg/dl  
(1 = true; 0 = false)',  
    ylabel='Count', width=500, height=450, legend_cols=2, legend_position='top_right'  
)
```

Out[14]:

```
In [15]: have_disease = data.loc[data['target']==1, 'restecg'].value_counts().hvplot.bar(alpha=0.4)
```

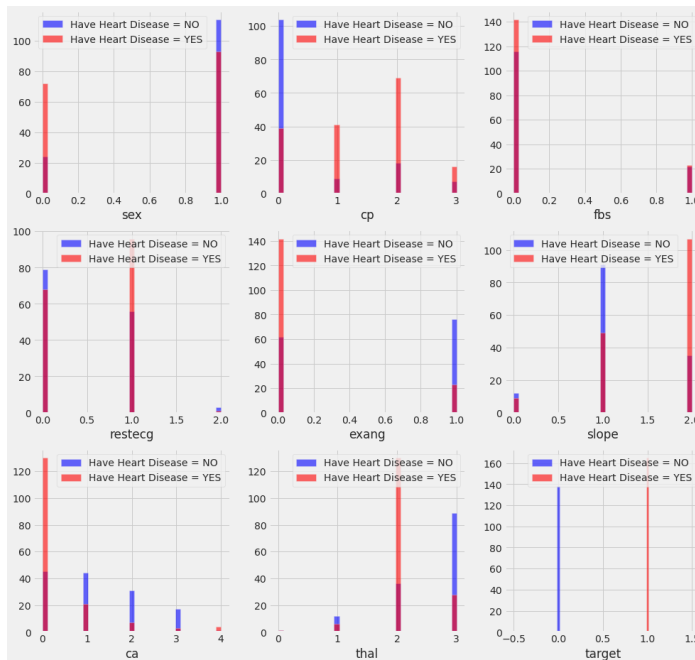
```
no_disease = data.loc[data['target']==0, 'restecg'].value_counts().hvplot.bar(alpha=0.4)

(no_disease * have_disease).opts(
    title="Heart Disease by resting electrocardiographic results", xlabel='resting electrocardiographic results',
    ylabel='Count', width=500, height=450, legend_cols=2, legend_position='top_right'
)
```

Out[15]:

In [16]: plt.figure(figsize=(15, 15))

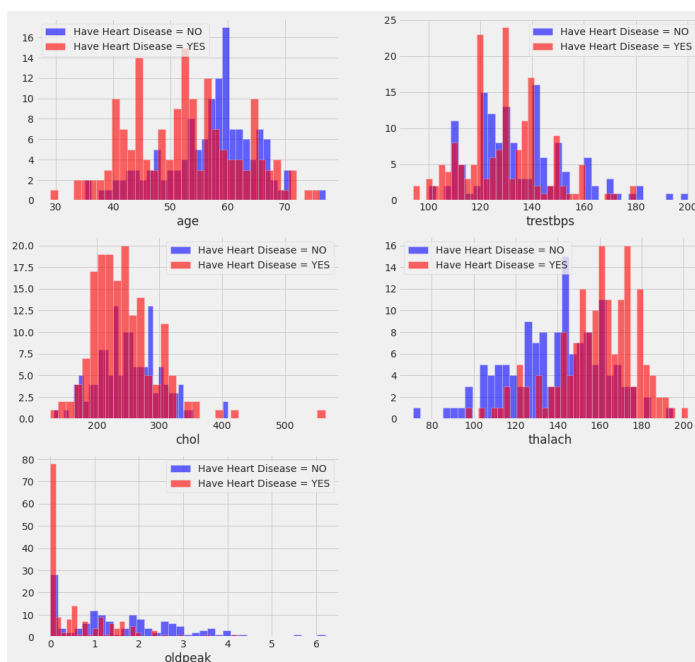
```
for i, column in enumerate(categorical_val, 1):
    plt.subplot(3, 3, i)
    data[data["target"] == 0][column].hist(bins=35, color='blue', label='Have Heart Disease = NO', alpha=0.6)
    data[data["target"] == 1][column].hist(bins=35, color='red', label='Have Heart Disease = YES', alpha=0.6)
    plt.legend()
    plt.xlabel(column)
```



In [17]: plt.figure(figsize=(15, 15))

```
for i, column in enumerate(continous_val, 1):
    plt.subplot(3, 2, i)
    data[data["target"] == 0][column].hist(bins=35, color='blue', label='Have Heart Disease = NO', alpha=0.6)
    data[data["target"] == 1][column].hist(bins=35, color='red', label='Have Heart Disease = YES', alpha=0.6)
    plt.legend()
```

```
plt.xlabel(column)
```




```
In [18]: # Create another figure
plt.figure(figsize=(9, 7))

# Scatter with positive examples
plt.scatter(data.age[data.target==1],
            data.thalach[data.target==1],
            c="salmon")


# Scatter with negative examples
plt.scatter(data.age[data.target==0],
            data.thalach[data.target==0],
            c="lightblue")

# Add some helpful info
plt.title("Heart Disease in function of Age and Max Heart Rate")
plt.xlabel("Age")
plt.ylabel("Max Heart Rate")
plt.legend(["Disease", "No Disease"]);
```



```
In [19]: # Let's make our correlation matrix a little
prettier
corr_matrix = data.corr()
fig, ax = plt.subplots(figsize=(15, 15))
ax = sns.heatmap(corr_matrix,
                  annot=True,
                  linewidths=0.5,
                  fmt=".2f",
                  cmap="YlGnBu");
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

Out[19]: (14.5, -0.5)



```
In [20]: categorical_val.remove('target')
dataset = pd.get_dummies(data, columns = cate
gorical_val)
```

```
In [21]: dataset.head()
```

```
Out[21]:
```

	age	trestbps	chol	thalach	oldpeak	target	sex
0	63	145	233	150	2.30	1	0
1	37	130	250	187	3.50	1	0
2	41	130	204	172	1.40	1	1
3	56	120	236	178	0.80	1	0
4	57	120	354	163	0.60	1	1

5 rows × 31 columns

```
In [22]: print(data.columns)
```

```
In [22]: print(data.columns)
print(dataset.columns)

Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
       'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
      dtype='object')
Index(['age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'target', 'sex_0',
       'sex_1', 'cp_0', 'cp_1', 'cp_2', 'cp_3', 'fbs_0', 'fbs_1', 'restecg_0',
       'restecg_1', 'restecg_2', 'exang_0', 'exang_1', 'slope_0', 'slope_1',
       'slope_2', 'ca_0', 'ca_1', 'ca_2', 'ca_3', 'ca_4', 'thal_0', 'thal_1',
       'thal_2', 'thal_3'],
      dtype='object')
```

```
In [23]: from sklearn.preprocessing import StandardScaler


s_sc = StandardScaler()
col_to_scale = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
dataset[col_to_scale] = s_sc.fit_transform(dataset[col_to_scale])
```

```
In [24]: dataset.head()
```

```
Out[24]:
```

	age	trestbps	chol	thalach	oldpeak	target	sex
0	0.95	0.76	-0.26	0.02	1.09	1	0
1	-1.92	-0.09	0.07	1.63	2.12	1	0
2	-1.47	-0.09	-0.82	0.98	0.31	1	1
3	0.18	-0.66	-0.20	1.24	-0.21	1	0
4	0.29	-0.66	2.08	0.58	-0.38	1	1

5 rows × 31 columns



```
In [25]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

def print_score(clf, X_train, y_train, X_test, y_test, train=True):
    if train:
        pred = clf.predict(X_train)
        clf_report = pd.DataFrame(classification_report(y_train, pred, output_dict=True))
        print("Train Result:\n")
        print(f"Accuracy Score: {accuracy_score(y_train, pred) * 100:.2f}%")
        print("")
        print(f"CLASSIFICATION REPORT:\n{clf_report}")
```

```
report}")
    print("")
    print(f"Confusion Matrix: \n {confusion_matrix(y_train, pred)}\n")

    elif train==False:
        pred = clf.predict(X_test)
        clf_report = pd.DataFrame(classification_report(y_test, pred, output_dict=True))
        print("Test Result:\n")
        print(f"Accuracy Score: {accuracy_score(y_test, pred) * 100:.2f}%")
        print("")
        print(f"CLASSIFICATION REPORT:\n{clf_report}")
        print("")
        print(f"Confusion Matrix: \n {confusion_matrix(y_test, pred)}\n")
```

In [26]: **from sklearn.model_selection import train_test_split**

```
X = dataset.drop('target', axis=1)
y = dataset.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

1.Logistic Regression

In [27]: **from sklearn.linear_model import LogisticRegression**

```
lr_clf = LogisticRegression(solver='liblinear')
lr_clf.fit(X_train, y_train)

print_score(lr_clf, X_train, y_train, X_test, y_test, train=True)
print_score(lr_clf, X_train, y_train, X_test, y_test, train=False)
```

Train Result:

Accuracy Score: 86.79%

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg
weighted avg				
precision	0.88	0.86	0.87	0.87
recall	0.82	0.90	0.87	0.86
f1-score	0.85	0.88	0.87	0.87


```
support    97.00 115.00      0.87    212.00
212.00
```

Confusion Matrix:

```
[[ 80 17]
 [ 11 104]]
```

Test Result:

Accuracy Score: 86.81%

CLASSIFICATION REPORT:

```

              0      1  accuracy  macro avg  w
weighted avg
precision    0.87  0.87      0.87      0.87
0.87
recall       0.83  0.90      0.87      0.86
0.87
f1-score     0.85  0.88      0.87      0.87
0.87
support     41.00 50.00      0.87     91.00
91.00
```

Confusion Matrix:

```
[[34  7]
 [ 5 45]]
```

```
In [28]: test_score = accuracy_score(y_test, lr_clf.pr
edict(X_test)) * 100
train_score = accuracy_score(y_train, lr_clf.
predict(X_train)) * 100

results_df = pd.DataFrame(data=[["Logistic Re
gression", train_score, test_score]],
                           columns=['Model',
'Training Accuracy %', 'Testing Accuracy %'])
results_df
```

Out[28]:

	Model	Training Accuracy %	Testing Accuracy %
0	Logistic Regression	86.79	86.81

2.Random Forest

```
In [29]: from sklearn.ensemble import RandomForestClas
sifier
from sklearn.model_selection import Randomize
dSearchCV

rf_clf = RandomForestClassifier(n_estimators=
1000, random state=42)
```

```
rf_clf.fit(X_train, y_train)

print_score(rf_clf, X_train, y_train, X_test,
y_test, train=True)
print_score(rf_clf, X_train, y_train, X_test,
y_test, train=False)
```

Train Result:

Accuracy Score: 100.00%

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg
weighted avg				
precision	1.00	1.00	1.00	1.00
1.00				
recall	1.00	1.00	1.00	1.00
1.00				
f1-score	1.00	1.00	1.00	1.00
1.00				
support	97.00	115.00	1.00	212.00
212.00				

Confusion Matrix:

```
[[ 97  0]
 [ 0 115]]
```

Test Result:

Accuracy Score: 82.42%

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	w
weighted avg					
precision	0.80	0.84	0.82	0.82	
0.82					
recall	0.80	0.84	0.82	0.82	
0.82					
f1-score	0.80	0.84	0.82	0.82	
0.82					
support	41.00	50.00	0.82	91.00	
91.00					

Confusion Matrix:

```
[[33  8]
 [ 8 42]]
```

In [30]:

```
test_score = accuracy_score(y_test, rf_clf.predict(X_test)) * 100
train_score = accuracy_score(y_train, rf_clf.predict(X_train)) * 100
```

```
results_df_2 = pd.DataFrame(data=[["Random Forest Classifier", train_score, test_score]],
                             columns=['Model',
                                     'Training Accuracy %', 'Testing Accuracy %'])
results_df = results_df.append(results_df_2)
```

```
results_df = results_df.append(results_df_2,
                                ignore_index=True)
results_df
```

Out[30]:

	Model	Training Accuracy %	Testing Accuracy %
0	Logistic Regression	86.79	86.81
1	Random Forest Classifier	100.00	82.42

Accuracy of Logistic Regression

```
In [31]: test_score = accuracy_score(y_test, lr_clf.predict(X_test)) * 100
train_score = accuracy_score(y_train, lr_clf.predict(X_train)) * 100

tuning_results_df = pd.DataFrame(data=[["Tuned Logistic Regression", train_score, test_score]],
                                columns=['Model',
                                          'Training Accuracy %', 'Testing Accuracy %'])
tuning_results_df
```

Out[31]:

	Model	Training Accuracy %	Testing Accuracy %
0	Tuned Logistic Regression	86.79	86.81

Accuracy of Random Forest

```
In [32]: test_score = accuracy_score(y_test, rf_clf.predict(X_test)) * 100
train_score = accuracy_score(y_train, rf_clf.predict(X_train)) * 100

results_df_2 = pd.DataFrame(data=[["Tuned Random Forest Classifier", train_score, test_score]],
                                columns=['Model',
                                          'Training Accuracy %', 'Testing Accuracy %'])
tuning_results_df = tuning_results_df.append(results_df_2, ignore_index=True)
tuning_results_df
```

Out[32]:

	Model	Training Accuracy %	Testing Accuracy %
--	-------	---------------------	--------------------

	Model	Training Accuracy %	Testing Accuracy %
0	Tuned Logistic Regression	86.79	86.81
1	Tuned Random Forest Classifier	100.00	82.42

Random forest Feature

```
In [33]: def feature_imp(df, model):
         fi = pd.DataFrame()
         fi["feature"] = df.columns
         fi["importance"] = model.feature_importances_
         return fi.sort_values(by="importance", ascending=False)
```

```
In [34]: feature_imp(X, rf_clf).plot(kind='barh', figsize=(12,7), legend=False)
```

Out[34]: <AxesSubplot:>

